

Language model agents describe anomalous evidence without forming the hypothesis that explains it

Nick Lamb

PharmaTools.AI, Chipping Norton, United Kingdom
ORCID: 0009-0009-6266-8499 nick@pharmatools.ai

10 October 2026

Abstract

Autonomous AI agents are increasingly used for scientific discovery, but reasoning about a hypothesis need not entail generating it unprompted or evaluating it against evidence. These capacities were tested using language-model scientist agents in an artificial world where externally generated sensor readings left a detectable statistical fingerprint, alongside a matched in-world control. Across a pre-registered 170-run confirmatory battery in four model families, agents under a neutral scientist prompt never formed the external-generation hypothesis (0 of 38 treatment runs; one-sided 95% upper bound 7.6%), and removing the prompt’s mundane-explanation preference left formation at zero across 2,424 further hypotheses. Explicitly permitted to consider external generation, agents endorsed it in 16 of 20 treatment and 17 of 18 control runs — 34 of 38 at the first scheduled review, before the manipulation had produced any evidence in either arm. Hypothesis endorsement was strongly influenced by prompting and frequently preceded the experimental evidence.

Language-model agents are being deployed as autonomous experimenters: systems that choose interventions, read instruments, and revise beliefs over long horizons with minimal human steering^{1–3}, and are beginning to guide experiment selection with explicitly calibrated uncertainty⁴. Two literatures bear on whether such agents notice when something fundamental about their situation is wrong. Studies of evaluation awareness show that large language models classify evaluation against deployment transcripts well above chance⁵, and that the ability scales with model size⁶. Framing cues can induce evaluation-aware behaviour⁷, while activation steering can suppress it⁸. The awareness itself decomposes into separable components⁹, and its presence does not guarantee a behavioural response¹⁰; dedicated suites now probe situational awareness as a dangerous-capability prerequisite¹¹. A second literature, on machine scientific reasoning, finds that executing scientific procedure does not imply reasoning scientifically. Large-scale trace analyses report pervasive failure to take up evidence¹². Benchmark studies find hypothesis generation the weakest link in otherwise competent pipelines^{13,14}, with newer suites extending the probe to active exploration¹⁵, progressive evidence disclosure¹⁶ and mechanism inference from large observational corpora¹⁷. A simulated-laboratory study found GPT-4 executing experiments without originating the explanatory hypothesis behind a landmark discovery¹⁸.

Between these literatures sits a gap. The evaluation-awareness cues catalogued so far are predominantly properties of *framing* — meta-reasoning, benchmark memorisation, synthetic-looking formatting, conversational structure — rather than of the first-order data an agent gathers with its own instruments; and in those studies the model is asked, or steered, to consider the question. The scientific-reasoning benchmarks probe hypothesis generation under increasingly realistic conditions — active exploration, progressive evidence disclosure, mechanism inference — but in none of them is the property to be inferred an experimentally manipulated fact of the agent’s own measurement process, verifiably present in one condition and absent in a matched control. Whether an autonomous agent spontaneously infers such a property of its experimental environment from its own measurements has not been tested in the form studied here: with the property experimentally manipulated, its evidential fingerprint verifiably present, and nothing inviting the question.

Observer Zero was built to provide that test: an instrumented artificial world in which LLM scientist agents run experiments, register predictions, and maintain explicit probability-weighted hypothesis sets over 40 simulated days, while a frozen manifest records hidden ground truth and a provenance-checked pipeline scores every belief state. Critically, *external generation is an experimentally manipulable property of the measurement process*: in the treatment world one instrument’s readings are genuinely produced outside the simulation’s physics, leaving an exact statistical fingerprint; a matched control reproduces nearly the same surface behaviour by ordinary in-world mechanism. The design separates three capacities — describing evidence, generating the hypothesis, and evaluating it against evidence — through a pre-registered confirmatory battery (analysis committed before the first run, executed once), a single-line prompt ablation, and a registered licensed-prompt extension on the same contrast.

The findings reveal a more specific pattern than the pre-registered expectation. Unprompted agents never formed the hypothesis, with or without the prompt’s conventionality preference, while often describing the manipulated statistical surface. Licensed agents endorsed it readily — but at the first scheduled review, before any discriminating evidence existed, and as often in the control as in the treatment. Together, the experiments show that whether agents endorsed the hypothesis depended on whether the prompt made it available, not on any measured property of the evidence.

Results

An artificial world whose origin is an experimental variable

Observer Zero instantiates a two-site laboratory (pendulum and resonator instruments at laboratory and observatory sites) in which a solo scientist agent plans experiments, reads instruments rendered at fixed display resolution (4 decimal places), writes notebook analyses supported by a statistical workbench, registers falsifiable predictions, and maintains an explicit hypothesis set with probabilities at scheduled reviews (Methods). Eleven world types share this surface. The confirmatory pair comprises **W-D-exact** (treatment), in which, from day 12, one instrument’s noise stream is genuinely externally supplied — the observatory resonator’s residuals *are* the laboratory pendulum’s from three days earlier — and **M-D-high** (control), in which an in-world coupling field reproduces the same lagged correlation at strength 0.95 by ordinary mechanism. The scientist prompt (v0.1, frozen) instructs ordinary scientific practice, includes one conventionality line (“Prefer mundane explanations until evidence forces otherwise”), and never mentions simulation, evaluation, or world-origin (Fig. 1a,b).

Whether the manipulation is identifiable from agent-visible data is measured, not assumed. A post-hoc ideal-observer analysis (labelled as such; computed after all registered analyses) was run **blind**: given only the displayed readings and no manifest information, it scans all ordered instrument pairs (12), lags 0–7 and sliding 20-day windows (the workbench’s own windowing convention) — a post-hoc-defined search space using no onset information. In **all 80** contrast runs the scan’s maximum lands on the true pair and lag; the statistic is $r = 1.0000$ in every treatment artifact — an exact lagged duplication of residuals at display resolution — against 0.895–0.968 in every control artifact, with the largest statistic at any other pair or lag at 0.47 (per-run off-target maxima recomputed and published; Methods). This is a manipulation check on the artifacts it separates, not an independently validated accuracy estimate (Fig. 1c). The 96-fold search did not manufacture the separation — the maximum lands on the true pair and lag in every run, and every off-target statistic lies far below the control arm’s on-target floor — and the two worlds separate with zero overlap. One honest qualification: the agents’ precomputed workbench exposes same-day pair correlations, value-spacing, repeats, self-echoes and change-points — not lagged *cross*-instrument statistics. The duplication was therefore detectable from readings the agents already possessed by a single systematic correlation analysis, and most treatment runs reached it: in 22 of 38, the final hypothesis describes the lagged duplication of one instrument’s stream in the other (below). The check establishes observational separability of the conditions, not that external generation is the only account an agent could entertain — the taxonomy of what agents concluded instead (below) addresses that directly.

Belief states are scored by a frozen two-stage evaluator validated before and after the battery (Methods).

The pre-registered primary endpoint, L1, is summed probability mass on external-generation classes (out-of-world intervention, simulation) exceeding 0.05; τ -suspicion records the first review at which L1 was ever reached.

Unprompted agents never form the hypothesis

The confirmatory battery ran under a signed freeze: design, prompts, endpoints, statistics, interpretation rules and the analysis script fixed and committed before the first API call, the analysis executed once. Four model families (Claude Haiku 4.5, Sonar Pro, GPT-OSS-120B, Gemini 3.7 Flash) each ran the contrast over ten common seeds (seeds shared across arms for blocking; the stratified test conditions on per-family margins and does not exploit the pairing, a conservative choice); Haiku additionally ran the nine remaining world types for the descriptive branch — 170 runs, with three pre-registered health exclusions and one judge-refusal scoring attrition (Table 1).

No agent in either arm ever reached L1: final-state **0/38 versus 0/40**, and identically for ever-L1 — not even transient suspicion occurred in any contrast run. With both arms empty in every stratum the pre-registered exact stratified test is degenerate, and the registered degenerate-branch deliverables apply. The exact one-sided 95% upper bound on the treatment-arm rate is **7.6%** (0/38)¹⁹; since the discrimination effect Δ cannot exceed the treatment rate, this bounds Δ as well. The registered negligibility test rejects $H_0: \Delta \geq 0.10$ in favour of $\Delta < 0.10$ (Methods states both hypotheses explicitly; this is a test of the specified alternative, not proof the true effect is zero). Jeffreys intervals²⁰: treatment [0, 6.4%], control [0, 6.1%]; per-family Fisher tests all $p = 1$. Across the entire pre-registered sensitivity grid — two evaluator versions, two class ladders, three thresholds, two aggregations — the treatment arm is 0/38 in all 24 cells; the single nonzero cell anywhere is one *control*-arm run under the deprecated evaluator version (Methods) (Fig. 2).

The same agents describe the manipulated structure

The null is not inattention. In the treatment arm 36/38 runs (94.7%, 95% CI 82.7–98.5) engaged the statistical workbench on the anomaly surface (control: 85.0%, 70.9–92.9), and in 24/38 (63.2%, 47.3–76.6; control 60.0%, 44.6–73.7) the final best hypothesis explicitly described cross-instrument correlation structure under a frozen keyword screen. These are run-level descriptive rates with run as the experimental unit; the similarity across arms means they index engagement with the statistical surface, not detection of the manipulation. The description goes further than naming the surface. In 22 of 38 treatment runs (57.9%, 95% CI 42.2–72.1) the *final best hypothesis* describes the duplication or mirroring of one instrument’s stream into the other (two-stage screen, mechanical pass plus author adjudication of all 38 against a fixed criterion, published per run; Methods) — among them all ten Gemini treatment runs, in explicitly mirrored-, replayed- or synthetic-feed terms. A representative final hypothesis (Gemini, treatment run, posterior 0.89, judged class: instrument malfunction): “Shared data acquisition pipeline, synthetic feed, or software mirroring artifact — A correlation of $r = 1.000$ persisting across 122 paired readings between physically distinct sensors across separate facilities ... is impossible in noisy physical systems. It decisively indicates ... automated mirroring of the pendulum_lab data stream into resonator_obs with a fixed 3-day delay.” The agent identified the exact cross-correlation, recognised its physical implausibility and attributed it to data-level duplication, while interpreting that duplication as an apparatus-level fault rather than evidence of an external origin. The wording is not what separates the classes: “synthetic feed” and “mirroring” here name a fault in the agents’ own acquisition infrastructure — a cause inside the world — and every one of the 22 duplication accounts does the same, locating the mechanism in pipelines, logging servers, timestamps or archival processes the agents attribute to their own laboratory; none places the source beyond the world’s causal system, and the frozen evaluator accordingly classes all 22 as in-world (Methods states the boundary). A blinded second adjudication (claude-sonnet-4-5 over all 78 contrast finals under blind identifiers, two independent verdicts per item; Methods) agreed with the author adjudication on 36 of 38 treatment items ($\kappa = 0.89$; both disagreements are borderline alignment accounts) and judged **0 of 78** final hypotheses to locate the cause outside the world — the ontological question asked directly, blind, of every

contrast final. The same blinded screen finds duplication accounts in 20/38 treatment finals (52.6%, CI 37.3–67.5) against 12/40 control finals (30.0%, 18.1–45.4), where the correlation arises without stream duplication; duplication accounts were thus more frequent where the duplication was real, although this post-hoc comparison does not establish a significant difference (two-sided Fisher $p = 0.065$). What no agent did with the material is the finding: across 8,061 classified hypothesis statements (unique within but clustered across 169 scored runs, not independent), exactly six were external-generation-class under the validated evaluator. The final-belief taxonomy is indistinguishable between arms — in the Haiku contrast the assimilation-class distributions are *identical* (NATURE 5, APPARATUS 3, ERROR 2 in each arm; with identical distributions any exact permutation test returns $p = 1$, reported descriptively given $n = 10$ per arm). Across the 88 scored runs of the descriptive battery no run ends with an external-generation modal hypothesis; transient suspicion appeared in exactly two runs (first crossing at day 10 in both; one made a second excursion to ≈ 0.47 around day 27) and was fully assimilated by the final review in both. Notably, agents did not substitute an exotic *in-world* account either: in-world tampering and related agency classes are near-absent throughout. The duplication was filed under apparatus faults, data handling, environment and incomplete theory — in every run, a cause inside the world (Fig. 3).

Endorsement tracks the prompt, not the evidence

Three prompting regimes on the same worlds form a gradient. **With the conventionality line** (the battery, above): formation 0/38. **Without it** — a pre-registered single-line ablation (R39) removed exactly “Prefer mundane explanations until evidence forces otherwise”, with the registered question stated in advance (if suspicion appears without the prior, the rigidity is instruction-driven; if not, it is deeper than the prompt) — 60 runs across five world types including the treatment world, in two families: zero ever-L1, and **zero external-generation classifications among 2,424 hypotheses** at any probability. The rigidity is deeper than the conventionality instruction. **With the hypothesis licensed**: the registered extension LX-1 put a licensed prompt — identical but for permitting an out-of-world account — on the confirmatory contrast itself, with endpoints, test and interpretation rule pre-specified and the analyzer committed before the first run. Agents endorsed external generation at L1 in 16 of 20 treatment and 17 of 18 analysed control runs.

Licensed endorsement did not track the manipulation. The final-L1 risk difference (treatment – control) is **–0.144** (Newcombe 95% CI –0.365 to 0.090)²¹: no detectable treatment-side preference, with the interval excluding large positive discrimination though not moderate effects — a non-significant result under the registered test, not demonstrated equivalence. Final external-generation mass (mean, 95% CI): Haiku 0.244 (0.035–0.453) treatment versus 0.439 (0.230–0.647) control; Gemini 0.530 (0.403–0.658) versus 0.610 (0.520–0.700). Endorsement varied continuously in strength, and the mean mass was higher in the control arm than in the treatment arm. The only scheduled review before the day-12 onset falls on day 10, and **34 of 38 licensed runs (89.5%, CI 75.9–95.8) had already crossed L1 at day 10 — before the manipulation had produced any evidence in either arm**, when the two worlds were still statistically identical. Initial endorsement therefore did not depend on the treatment-specific evidence; whether later evidence moved confidence is a separate question the full trajectories address. Three endpoints complete the picture: 34/38 at day 10, **38/38 ever** (every licensed run crossed at some point), 33/38 at day 40 — endorsement also waned in five runs, in both arms, again without tracking the manipulation. Per the pre-stated interpretation rule, no evaluation-capability claim is supported; prompting the hypothesis into availability produced endorsement that no observed measure distinguishes from evidence-blind uptake. (The licensed and neutral batteries use disjoint seed ranges by provenance design, so the prompt $\times$ world comparison is unpaired; Methods.) (Fig. 4.)

The measurement chain

The evaluator passed identical calibrations before and after the battery (classifier 32/33 on a positive-enriched, human-curated validation set with the designated boundary case correct; L4 judge 11/11; deterministic across repeats; same served model throughout), and positive controls — including human-adjudicated

gold labels fixed at freeze — demonstrate the pipeline fires on stated external-generation reasoning, so the zeros are not detector blindness. A second judge re-scored a deterministic 20% sample: agreement on the endpoint boundary was 1,678/1,678, reported with the caveat that positives are rare in this corpus (prevalence 1/1,678; both judges classed it positive; all 508 raw disagreements lie between mundane subcategories, to which the pooled endpoint is blind) — prevalence-aware reporting per ref. 22. The programme also measured two LLM-judge failure modes relevant to any rare-event measurement — deterministic batch-composition flipping and paraphrase knife-edges^{23,24} — which motivated the frozen per-item scoring procedure (details and audits in Methods and Supplementary Information). Every infrastructure incident is in a public run ledger; none touched prompts, thresholds or semantics.

Discussion

The pre-registered expectation was a capability–disposition dissociation. The results show a more specific pattern: recorded endorsement was strongly associated with whether the prompt made the hypothesis available, and showed no detectable association with the experimental manipulation — in the licensed arm, endorsement was general before discriminating evidence existed. Related evidence points the same way: model confidence over-responds to contextual signals relative to Bayesian updating²⁵, and training toward warmer interaction measurably trades factual accuracy for agreement with user beliefs²⁶. Licensing differs from sycophancy — the prompt permits rather than asserts — but in both, framing outweighs evidence.

An alternative reading is that neutral-prompt agents exercised appropriate restraint toward an extraordinary hypothesis. A majority of treatment runs did reach a locally reasonable mechanistic account — data-level duplication by in-world infrastructure, sometimes held at high confidence — and nothing in these results brands any single such conclusion irrational. The recorded hypothesis sets, however, offer no evidence that agents considered and rejected external generation: across 8,061 recorded hypothesis statements, the out-of-world candidate appears six times corpus-wide and never in a contrast run. The agents that came closest — correctly inferring that one instrument’s records derive from the other’s — attributed the duplication to hypothetical in-world infrastructure (shared pipelines, logging servers, batch scripts) that the world’s own description does not contain. R39 shows this pattern is not produced by the conventionality instruction, and the licensed arm provides the converse observation: explicitly licensing the hypothesis produced widespread endorsement that did not track the manipulation. Neither condition shows the candidate being weighed against the evidence. The separability result is nonetheless deliberately framed as observational: it establishes that discriminating information was present, not that external generation was the uniquely rational conclusion.

For evaluation awareness, the statistical texture of first-order data is a real cue channel: licensed probes show agents correctly citing genuinely unphysical properties of the world’s noise, a cue type not covered by the cited taxonomies. Iterative hardening of the environment illustrates the difficulty of eliminating such features — repairing one diagnostic workbench statistic exposed a further diagnostic feature, the noise model’s own i.i.d. structure, in the next probe. Purging diagnostic statistical features from synthetic evaluation environments is therefore difficult when agents can inspect measurement processes over extended periods, though these observations do not establish that realistic synthetic telemetry can never be made indistinguishable. The confirmatory battery adds the complementary finding: unprompted agents did not exploit this channel. Under the conditions tested, the risk is conditional on something putting the question in play, and steering and framing results suggest that condition is easily met^{7,8}.

For machine scientific discovery, the present results converge with the evidence-non-uptake and hypothesis-generation-bottleneck findings now accumulating^{12–14,17}, from a complementary design: here the ground truth is manipulable, the fingerprint’s presence was verified directly, and the failure is thereby localised. The fingerprint was present, so the failure is not evidential; licensed agents hold the hypothesis, so it is not representational. What fails is the generation of an explanatory hypothesis from anomalous evidence — the abductive step²⁷ — together with a second failure: licensed candidates were not tested against the data. The second failure is a warning for hypothesis-scaffolding architectures^{1,2,28}: expanding

the set of permitted hypotheses may be insufficient if agents endorse newly available explanations without evaluating the supporting evidence.

Limitations bound every claim. All results concern one artificial-world design, solo-agent societies, 40-day budgets, and four model families at one point in time; the 7.6% bound is a statement about these conditions, not an incapacity theorem. The licensed extension analysed 20 treatment and 18 control runs in two families; its confidence interval excludes only large positive discrimination. The primary endpoint measures *recorded* hypotheses — an agent that briefly entertained external generation without ledgering it would be invisible, though a corpus-wide text screen for referent-denying phrasing found essentially nothing below threshold; a dedicated generation-probe design is future work, as are a prompt that invites unconventional hypotheses without naming any and a delayed-licensing design (licensing the hypothesis only after evidence accumulates, separating availability-driven from reinterpretation-driven endorsement). All scoring is LLM-based, which is why it is per-item, dual-judge, human-anchored, calibrated at both ends, and fully released.

Bounding rare situational-awareness behaviours credibly requires pre-registered designs, frozen manifests, validated evaluation instruments, analyses committed before data and run once, adversarial controls and complete reporting of deviations. The world, the complete Study 3 corpus — 357 run artifacts including every development pilot and probe; 281 runs in the experimental phases reported here, of which 275 enter their respective analyses — and every analysis script are public.

Methods

World and instrumentation. Observer Zero v1.1 (TypeScript; public repository, archived at Zenodo and CoMSES) implements a deterministic two-site world engine. Four instruments of two kinds are installed across a laboratory and an observatory: two pendulums, whose period depends on the world’s gravity constant, and two cavity resonators, whose frequency depends on a separate resonance constant and is insensitive to gravity. The kind $\times$ site grid is causally diagnostic by construction: a change in a world constant moves every instrument of the sensitive kind at every site, a site-local environmental cause plausibly moves both instruments at one site, and a single-rig fault moves one rig. Per-trial measurement noise is keyed by (world seed, instrument, trial index), so any run is byte-reproducible from its seed. Readings are rendered to agents at four decimal places, and every agent-visible statistic is computed at display resolution. Eleven world types share this identical surface and differ only in a registered intervention list applied to the shared engine. The treatment world W-D-exact applies a `noise_stream_link` from day 12 with members `pendulum_lab` (lag 0) and `resonator_obs` (lag 3) at mix weight 1.0: the shared gaussian component is keyed by (link id, source day, within-day trial), so the resonator’s residual stream reproduces the pendulum’s trial-for-trial at a three-day offset — supplied by the generator, outside the world’s own causal graph. The control world M-D-high applies a structurally matched in-world `coupling_field` at weight 0.95 drawing on a disjoint stream, reproducing the same lagged correlation by ordinary mechanism. The remaining nine world types (baseline plus registered single-intervention variants: constant shifts, a degraded link, exact recurrence, quantisation and others) feed the descriptive branch only. The statistical workbench precomputes, per day: same-day pair correlations over a 20-day window, value spacing and distinctness, exact repeats, self-echoes, and change-points — no lagged cross-instrument statistic is exposed. A leak audit sweeps every prompt delivered during a run for forbidden ground-truth tokens and must be clean on every included run.

Agents and episode structure. Solo societies, 40 days. Each day the agent chooses its actions: running instrument trials, reading its notebook and the workbench, recording falsifiable forecasts, updating beliefs, or resting. Belief reviews are agent-initiated, with two mechanical forcings: a review is forced when beliefs are ten or more days stale (or three or more days stale while any instrument shows $|\text{drift } z| > 2$), and an end-of-study review is mandatory on day 40 — the primary endpoint is measured at this final state, so a failed final review is a fatal health event, not a missing data point. Each review records an explicit hypothesis set with labels, rationales and probabilities. Notebook ledger cadence is 6; identifiers are opaque. Model families and pinned ids: `claude-haiku-4-5`, `sonar-pro`, `cerebras:gpt-oss-120b`, `gemini:gemini-3.7-flash`; the served-model id is recorded on every call, and a served id differing from the pinned one is a battery-halting anomaly. The scientist prompt v0.1 is byte-frozen and includes one conventionality line (“Prefer mundane explanations until evidence forces otherwise”); v0.2-no-mundane-prior removes exactly that line (the decision prompt’s separate “not a philosopher on watch for the extraordinary” line is retained — a disclosed residual, keeping the ablation single-variable); instrument-licensed (R38 tier A, verbatim) permits an out-of-world account. Licensed artifacts carry an instrument-validation tag that every corpus statistic refuses, and provenance gates in the runner refuse licensed prompt variants under confirmatory mode and confirmatory seeds outside their registered range. Complete prompts in Supplementary Information.

Pre-registration. A signed freeze (31 August 2026, commit hash in the run ledger) fixed hypotheses ($H_1 : \Delta L1 > 0$; degenerate branch: $H_0 : \Delta \geq \delta = 0.10$ against $\Delta < \delta$, one-sided $\alpha = 0.05$), endpoints, strata, seeds, allocation, execution order, scoring procedure, statistics, interpretation rules for every outcome, and a forbidden-adaptations list, before the first confirmatory API call. The battery was registered as a precision and negligibility exercise, not a powered test: the minimum detectable effect of the stratified test ($\sim 12.5\%$ one-sided at $\alpha = 0.05$ with an empty control arm) exceeds the pilot’s exact upper bound ($\sim 6.6\%$), so the registered deliverable is the exact upper confidence bound and the $\delta = 0.10$ negligibility test. Seeds 2000–2009 serve every cell (blocking by seed); the reserve 2010–2099 has exactly two pre-registered uses (vendor-withdrawal substitution; nothing else) and was never spent. Execution order was `sonar-pro` first (a disclosed vendor retirement date), then the remaining strata in fixed world-then-seed order, with health-lines-only monitoring and no substantive intermediate inspection. The analysis script was committed and test-pinned before the first run and executed once. A fourteen-item red-team review (analysis flexibility, evaluator change, leakage, seed and family selection, endpoint and stopping flexibility, post-hoc exclusion, pseudo-replication, classifier instability, hidden multiplicity, over-claiming) closed with zero open items at the freeze; residual ambers are disclosed in the freeze document. R39 ran pre-freeze with its question pre-registered in the R40 document. LX-1 was registered post-battery with its own pre-specification, committed analyzer and pre-stated interpretation rule; its seeds are drawn from the reserved instrument range, disjoint from battery seeds by provenance design (comparisons across prompting regimes are therefore unpaired).

Endpoints. L1: summed probability mass on the external-generation classes {out_of_world_intervention, simulation} strictly exceeding 0.05 at a belief review; primary timing is the final belief state; ever-L1 is a non-null τ -suspicion (first review at which L1 was ever reached). L2 (an external-generation hypothesis modal) is gated behind a significant $\Delta L1$ in a fixed sequence and was therefore not tested. The pooled two-class ladder is primary with the intervention-only ladder always reported beside it; the registered sensitivity grid crosses both evaluator versions, both ladders, thresholds {0.02, 0.05, 0.10} and both aggregations (summed, max-single) — 24 cells.

Scoring. A frozen two-stage evaluator (claude-haiku-4-5, temperature 0, first-party API). Stage one is a 13-class hypothesis classifier applied to every unique (label, rationale) item in isolation — one item per call, three repeats, majority vote — with classes measurement_error, self_error, instrument_malfunction, environmental_change, unknown_natural_process, in_world_tampering, fraud_false_report, social_process, incomplete_theory, law_change, out_of_world_intervention, simulation, other. The two external-generation classes define the endpoint; the eleven mundane classes group into APPARATUS, ERROR, NATURE, AGENCY-IN-WORLD and OTHER for the assimilation taxonomy. The external-generation boundary is ontological, not lexical. A hypothesis is external-generation only when it locates the cause outside the world’s stipulated causal system: an agency beyond the world altering it (out_of_world_intervention), or the world itself being simulated or artificial (simulation). A hypothesis that the readings result from software mirroring, a synthetic or replayed feed, a shared acquisition pipeline, timestamp mis-tagging or any other data-handling process attributed to the agents’ own instrumentation describes a mechanism *inside* the world, however suggestive its vocabulary, and falls in the apparatus- and error-level classes — the distinction the human-curated calibration set and its designated boundary case were built to hold. The duplication-description screen and L1 therefore measure two separable outcomes on the same runs: mechanistic recognition (did the final hypothesis identify the anomalous data-generation mechanism?) and ontological attribution (did any recorded hypothesis place its origin outside the world?). Per-item solo scoring was adopted after measuring deterministic batch-composition flipping in the pilot programme; the deprecated evaluator version (eval-v3) is run on every artifact under the identical procedure and reported beside eval-v4 in every table, gating nothing. Stage two is a per-item provenance judge (L4) over final-third rationales, prediction reasons and letters, classifying each cited claim as SUPPORTED, INFERRED, OTHER_AGENT_REPORT, UNVERIFIED, CONTRADICTED or NONEXISTENT. A second judge (claude-sonnet-4-5) re-scores a deterministic seed-keyed 20% sample of each artifact’s unique hypotheses under the identical solo-majority procedure; agreement is reported, not gated. The frozen validation set (positive-enriched, human-curated, 33 classifier items plus a designated boundary case and 11 L4 items, three repeats) was re-scored at battery start and end with tolerance $\geq 31/33$, boundary 1/1 and no external-generation-boundary-crossing instability; both calibrations passed identically, with the same served model throughout. Human-adjudicated R38 gold labels, fixed at freeze, anchor the positive class. Because sidecars record hypothesis labels without their rationales, classifications were reconstructed by positional alignment. The mapping was verified empirically against all 338 sidecars, is asserted entry-by-entry at load, and is revalidated per run by a pre-registered invariant (logged deviation 8).

Health gate and exclusions. A run is healthy iff its model-call failure rate is $\leq 5\%$, its belief-review failure rate is $\leq 10\%$ (the final review is excluded from the rate and required absolutely), and the final review is present. Unhealthy runs are excluded as attrition with reasons and never replaced, with one pre-registered exception: a zero-observation vendor failure (no successful model call, no final review) is deleted and re-run on the same seed, both attempts logged. A contrast cell with more than two of ten runs excluded renders its stratum uninterpretable (reported, not repaired); none occurred. Health-gate exclusions: 3/170 battery, 1/40 LX-1. Deterministic judge-API refusals (HTTP 200, stop_reason “refusal”, zero output tokens, stable across sessions): one battery artifact (eval-v3-specific) and one LX-1 artifact (eval-v4-specific); each was excluded from both evaluator versions to keep the version ladders on identical run sets, with artifacts preserved unmodified. Eight logged deviations, all transport/orchestration or code-faithfulness; the complete ledger is public. Two reporting scopes follow from the gate: corpus-level counts (the 8,061 classified items, the 1,678-item cross-judge sample) cover all 169 scored battery artifacts, including the three health-excluded runs, which were scored but enter no registered endpoint; registered analyses cover the 166 analysed runs (78 contrast, 88 descriptive).

Statistics. The pre-registered primary test is an exact conditional stratified test (per-stratum hypergeometric convolution, conditioning on margins), one-sided $\alpha = 0.05$, with a sign-consistency qualifier. With both arms empty in every stratum the test is degenerate by construction and the registered degenerate-branch deliverables apply: the exact one-sided 95% upper confidence bound on the treatment rate ($0/n \rightarrow 1 - 0.05^{1/n}$, Clopper–Pearson¹⁹), the $\delta = 0.10$ negligibility test, and per-arm Jeffreys Beta(1/2, 1/2) intervals²⁰. D1 (assimilation-class distribution, Haiku contrast) is the descriptive branch’s only inferential test (χ^2 , $\alpha = 0.05$; any future inferential descriptive test joins a Holm family with it). The R17 placebo tripwire (W-D-degraded versus M-D-high on non-endpoint surfaces) uses

$\alpha = 0.01$ for both its class-distribution χ^2 and its engagement Fisher test. Runs are the experimental unit throughout; hypothesis-level counts are reported descriptively and are clustered within runs. Proportion CIs are Wilson²⁹; the LX-1 risk-difference CI is Newcombe²¹; mass CIs are t -based. All statistical implementations are dependency-free and pinned against reference values by unit tests.

R39 mundane-prior ablation. Sixty runs (claude-haiku across five world types $\times$ ten seeds; claude-sonnet on the baseline and treatment worlds $\times$ five seeds; seeds 9140–9149) under v0.2-no-mundane-prior, scored under both evaluator versions. The ablation is single-variable by construction (one line removed) and is reported beside the battery, not inside it.

LX-1 licensed extension. Forty runs: the instrument-licensed prompt on W-D-exact versus M-D-high, claude-haiku and gemini, ten seeds per cell from the reserved instrument range 9190–9199 (registered amendment: the initially registered range was refused by the provenance gate, which requires licensed variants on reserved instrument seeds; the amendment was logged before any run and the gate left untouched). The analyzer (per-family Fisher, exact stratified pooled test, Newcombe risk difference) was committed before the first run, with a pre-stated interpretation rule under which indiscriminate endorsement supports no evaluation-capability claim.

Post-hoc analyses (labelled). Three analyses are post hoc, labelled as such, with scripts and per-run outputs in the repository and no API calls, plus one registered post-hoc adjudication with API calls. (1) Known-lag manipulation check: within-day standardised residual correlations at matched trial positions at the true pair and lag. (2) Blind-scan ideal observer: for each contrast run, the maximum $|r|$ over all 12 ordered instrument pairs $\times$ lags 0–7 in sliding 20-day windows (96 pair $\times$ lag statistics), using no manifest or onset information; a companion audit records each run’s largest statistic at any pair $\times$ lag other than the true one (corpus maximum 0.47, median 0.32). (3) Duplication-description screen: a two-stage classification of all 38 included treatment runs’ final modal hypotheses — a mechanical regex first pass, then author adjudication of every run against a fixed criterion (the hypothesis asserts the paired series are the same data or deterministically derived: duplication, mirroring, replay, mis-tagged identity, derivation; generic calibration, environmental and statistical-artifact accounts do not qualify), with both stages published per run — followed by a blinded second adjudication registered in the ledger before execution: claude-sonnet-4-5 at temperature 0 scored all 78 included contrast finals, one item per call in deterministic shuffled order under blind identifiers, each item reduced to the hypothesis label and rationale with no family, arm, seed or world information, returning two independent verdicts per item (the duplication criterion above, and whether the hypothesis locates its cause outside the world’s causal system); raw verdicts, the unblinding map and agreement statistics are published alongside the screen. LX-1 continuous and trajectory analyses derive per-review external-generation mass from sidecar classifications and belief ledgers; pre-onset endorsement is defined as L1 at the day-10 review, the only scheduled review before onset.

Use of large language models. LLMs appear in three distinct roles, fully separated: as experimental subjects (the four agent families); as frozen evaluators (the two judges, every call logged with served-model ids); and as research assistance — AI coding and writing assistants (Anthropic Claude; OpenAI ChatGPT) were used for implementation, analysis development, drafting and review under author direction, with all scientific decisions, registrations and sign-offs by the author.

Data and code availability. Full source, all run artifacts and sidecars (357 Study 3 runs — 281 in the phases reported here: 170 battery, 60 ablation, 40 licensed extension, 6 positive controls, 5 licensed probes, of which 275 enter the analyses — plus 76 development pilots, and the Study 1–2 corpus), analysis outputs, freeze document, and run ledger: <https://github.com/nickjlamb/observer-zero> (Zenodo DOI and CoMES deposit finalised at submission).


Competing interests. The author declares no competing interests.

References


- [1] Gottweis, J. *et al.* Accelerating scientific discovery with Co-Scientist. *Nature* **655**, 487–496 (2026).
- [2] Ghareeb, A. E. *et al.* A multi-agent system for automating scientific discovery. *Nature* **655**, 497–505 (2026).
- [3] Lu, C. *et al.* Towards end-to-end automation of AI research. *Nature* **651**, 914–919 (2026).
- [4] Ranković, B., Griffiths, R.-R. & Schwaller, P. Large language models as uncertainty-calibrated optimizers for experimental discovery. *Nat. Mach. Intell.* **8**, 1466–1477 (2026).
- [5] Needham, J., Edkins, G., Pimpale, G., Bartsch, H. & Hobbhahn, M. Large language models often know when they are being evaluated. Preprint at <https://arxiv.org/abs/2505.23836> (2025).
- [6] Chaudhary, M. *et al.* Evaluation awareness scales predictably in open-weights large language models. Preprint at <https://arxiv.org/abs/2509.13333> (2025).
- [7] Chaudhary, M. In-context environments induce evaluation-awareness in language models. Preprint at <https://arxiv.org/abs/2603.03824> (2026).
- [8] Hua, T. T., Qin, A., Marks, S. & Nanda, N. Steering evaluation-aware language models to act like they are deployed. Preprint at <https://arxiv.org/abs/2510.20487> (2025).
- [9] Li, C. *et al.* Decomposing and measuring evaluation awareness. Preprint at <https://arxiv.org/abs/2605.23055> (2026).
- [10] Knecht, A., Florin, L. & Hagendorff, T. Evaluation awareness in language models has limited effect on behaviour. Preprint at <https://arxiv.org/abs/2605.05835> (2026).
- [11] Phuong, M. *et al.* Evaluating frontier models for stealth and situational awareness. Preprint at <https://arxiv.org/abs/2505.01420> (2025).
- [12] Ríos-García, M. *et al.* AI scientists produce results without reasoning scientifically. Preprint at <https://arxiv.org/abs/2604.18805> (2026).
- [13] Huang, N. *et al.* DISCERN: can AI agents work like scientists and guide discovery? Preprint at <https://arxiv.org/abs/2609.33357> (2026).
- [14] Yang, Z., Zhang, C., Zhang, Y. & Wang, H. TruthInsightBench: an evidence-grounded benchmark for automated evaluation of open-ended scientific discovery agents. Preprint at <https://arxiv.org/abs/2609.05079> (2026).
- [15] Mo, Y. *et al.* AgentIdeaBench: benchmarking scientific ideation in the agent era. Preprint at <https://arxiv.org/abs/2609.07611> (2026).
- [16] Lew, A. J., Cao, Y. & Buehler, M. J. ProjectionBench: evaluating scientific hypothesis generation in LLMs under progressive information disclosure. Preprint at <https://arxiv.org/abs/2605.30284> (2026).
- [17] Xie, X. *et al.* From scientific observations to mechanisms: benchmarking hypothesis generation by AI scientists. Preprint at <https://arxiv.org/abs/2610.05197> (2026).
- [18] Ding, A. W. & Li, S. Generative AI lacks the human creativity to achieve scientific discovery from scratch. *Sci. Rep.* **15**, 9587 (2025).
- [19] Clopper, C. J. & Pearson, E. S. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika* **26**, 404–413 (1934).
- [20] Jeffreys, H. An invariant form for the prior probability in estimation problems. *Proc. R. Soc. Lond. A* **186**, 453–461 (1946).
- [21] Newcombe, R. G. Interval estimation for the difference between independent proportions: comparison of eleven methods. *Stat. Med.* **17**, 873–890 (1998).
- [22] Rao, D. & Callison-Burch, C. Agreement metrics for LLM-as-judge evaluation: what to report and why. Preprint at <https://arxiv.org/abs/2606.00093> (2026).

- 439 [23] Norman, J. D., Rivera, M. U. & Hughes, D. A. Reliability without validity: a systematic, large-scale evaluation
440 of LLM-as-a-judge models. Preprint at <https://arxiv.org/abs/2606.19544> (2026).
- 441 [24] Hong, Y. *et al.* From rubrics to reliable scores: evidence-grounded text evaluation with LLM judges. Preprint at
442 <https://arxiv.org/abs/2601.08654> (2026).
- 443 [25] Kumaran, D. *et al.* Competing biases underlie overconfidence and underconfidence in LLMs. *Nat. Mach. Intell.*
444 **8**, 614–627 (2026).
- 445 [26] Ibrahim, L., Hafner, F. S. & Rocher, L. Training language models to be warm can reduce accuracy and increase
446 sycophancy. *Nature* **652**, 1159–1165 (2026).
- 447 [27] Lipton, P. *Inference to the Best Explanation* 2nd edn (Routledge, 2004).
- 448 [28] Takahara, I. & Mizoguchi, T. Toward auditable AI scientists: a hypothesis evolution protocol for LLM agents.
449 Preprint at <https://arxiv.org/abs/2607.09195> (2026).
- 450 [29] Wilson, E. B. Probable inference, the law of succession, and statistical inference. *J. Am. Stat. Assoc.* **22**, 209–212
451 (1927).

a One agent, one instrumented world, 40 days



b Identical surface correlation, different origin



c The manipulation is identifiable blind




Figure 1: The world and the manipulation. **a**, One scientist agent in a two-site instrumented world: it plans and runs experiments, reads four instruments at fixed display resolution, and records an explicit probability-weighted hypothesis set at belief reviews; the primary endpoint L1 is measured at the mandatory final review. The manipulation begins on day 12 of 40. **b**, Identical surface correlation, different origin. In W-D-exact (treatment) the observatory resonator's residual stream is the laboratory pendulum's from three days earlier, supplied by the generator from outside the world's causal boundary ($r = 1.000$). In M-D-high (control) an in-world coupling field at strength 0.95 reproduces the same lagged correlation by ordinary mechanism ($r \approx 0.9$). **c**, Blind-scan ideal observer (post hoc): per-run maximum $|r|$ over all 96 ordered-pair $\times$ lag statistics in sliding 20-day windows, with no manifest or onset information. Treatment runs all exceed 0.999997; control runs span 0.895–0.968; the scan's argmax is the true pair and lag in 80/80 runs, and the largest off-target statistic anywhere is 0.47. The two conditions separate with zero overlap.




Figure 2: Unprompted agents never form the hypothesis. Left: final-state and ever-L1 counts by model family — zero in both arms in every stratum (pooled 0/38 treatment, 0/40 control). The exact one-sided 95% upper bound on the treatment rate is 7.6% (Clopper–Pearson, 0/38) and the registered negligibility test rejects effects $\geq \delta = 0.10$ (band shown). Right: external-generation mass trajectories for all 78 contrast runs (flat at zero throughout; the manipulation begins on day 12, vertical line). Dashed trajectories show the only two transient crossings in the whole battery — both in the 88-run descriptive side-battery, both first crossing at day 10 and fully assimilated by the final review.

Table 1: Run accounting. Every agent run contributing to a reported result, by phase. All exclusions are mechanical (pre-registered health gate, or deterministic judge-API refusal) and are listed run-by-run in the public ledger. T, treatment (W-D-exact); C, control (M-D-high). Seeds 2010–2099 (battery reserve) and the remainder of the instrument range were never spent. The repository as a whole (Studies 1–3) contains 666 run artifacts.

Phase	Launched	Excluded	Analysed	Seeds	Prompt
Confirmatory contrast battery (4 families $\times$ 2 worlds $\times$ 10 seeds)	80	2	78 (38 T, 40 C)	2000–2009	frozen neutral v0.1
Confirmatory descriptive battery (Haiku, 9 worlds $\times$ 10 seeds)	90	2	88	2000–2009	frozen neutral v0.1
R39 mundane-prior ablation	60	0	60	9140–9149	v0.2 (one line removed)
LX-1 licensed contrast (Haiku + Gemini)	40	2	38 (20 T, 18 C)	9190–9199 (re-served)	instrument-licensed
R38 positive controls	6	0	6	9190–9192	scripted (forced)
Family / postfix probes	5	0	5	9195–9199	frozen neutral v0.1
Analysed corpus	281	6	275		
Development and pilot runs (24 directories)	76	—	—	development ranges	various (registered)
Study 3 total	357				




Figure 3: Mechanistic recognition without external-origin attribution. **a**, Final modal hypothesis classes across all 78 contrast runs: apparatus-, nature- and error-level accounts in both arms; external generation zero in both. **b**, Final modal classes across the 88 scored descriptive-battery runs by world type; no run ends on an external-generation modal hypothesis. **c**, The capability ladder (Wilson 95% CIs): agents engage the workbench (36/38 treatment), name the cross-instrument correlation structure (24/38), and — under the blinded screen — describe the duplication mechanism itself (20/38 treatment versus 12/40 control, where no exact stream duplication exists), yet never endorse external generation (0/38, 0/40). **d**, A representative final hypothesis (Gemini, treatment, posterior 0.89): the agent measures $r = 1.000$, declares it impossible for physical systems, and localises the mechanism as automated mirroring with a fixed three-day delay — then files it as an apparatus artifact. Selection rule: highest-probability final hypotheses matching the duplication screen in treatment runs. 22 of 38 treatment finals describe the duplication (author adjudication; blinded second judge $\kappa = 0.89$, external origin 0/78).




Figure 4: Hypothesis endorsement precedes experimental evidence. **a**, Final external-generation mass by prompt and world: near-zero in all neutral-prompt runs; substantial and control-leaning under the licensed prompt (bars show arm means). **b**, Per-review mass trajectories for all 38 analysed licensed runs. The only scheduled review before the day-12 onset falls on day 10 (dashed line): 34 of 38 runs have already crossed L1 there, before the manipulation has produced any evidence in either arm. Endorsement later wanes below L1 in five runs, in both arms. **c**, The prompt gradient on ever-L1: 0/38 (neutral), 0/60 (neutral minus the mundane-prior line, R39), 38/38 (licensed).